

5G RRC Signaling Storm Mitigation

via a Proactive Digital Twin Architecture
with Continuous Calibration



Ömer Dikyol

Department of Computer Engineering
Istanbul Technical University



Outline

01

Problem Definition & Research Gap

Why availability matters; the storm; control-loop latency

02

System Model

MMPP/M/c/K queue and gNB RACH parameters

03

Proposed Method

Calibrated Digital Twin with EMA + asymmetric step rule

04

Results & Evaluation

Mirroring error, queue stability, delay CDF, ϵ -sensitivity

05

Conclusion

Findings, O-RAN deployment, future work

Why 5G availability matters

5G NR underpins mission-critical services — every device starts with an RRC connection request through the gNB RACH.

<10 ms

Round-trip latency required by autonomous vehicles

5×9

Reliability target for remote surgery and industrial IoT

100s–1000s

UEs sharing a single gNB simultaneously

gNB RACH at a glance

c = 54

preamble channels

K = 200

buffer capacity

Q ≥ K

requests dropped immediately



The RACH is a shared, limited resource with no built-in overload anticipation. RACH failure ⇒ no data plane, all services blocked.

What is an RRC signaling storm?

RRC and the RACH


RRC is the 3GPP Layer-3 protocol that manages connection setup. No data flows without a completed RRC connection.

How a storm forms

- Botnet floods Msg1 simultaneously, overwhelming $c = 54$ preamble slots
- Preamble collision $\Rightarrow$ both UEs fail $\Rightarrow$ retry
- gNB buffers pending Msg3 (RRC Requests) up to $K = 200$
- At $\lambda_{\text{high}} = 500$ req/s, buffer fills in 0.4 s
- Dropped devices retry — load never subsides

Net effect: RACH becomes a denial-of-service point — legitimate UEs cannot complete connection setup.

RACH procedure (Msg1–Msg4)



The RRC signaling storm attack


Attack mechanism

- Adversary coordinates a botnet of compromised devices
- Each device sends RRC connection requests at high rate
- Combined rate: $\lambda_{\text{high}} = 500 \text{ req/s}$ vs. normal $\lambda_{\text{low}} = 10 \text{ req/s}$ (50×)
- RACH buffer saturates within seconds
- Access requests are dropped — QoS failure

Termed “RRC Signaling Storm” in the literature [b1].

Impact: All connected services lose access until the attack subsides.

Attack flow



PROBLEM

Control-loop latency in ACB/UAC

Current defence: Access Class Barring (ACB) — 3GPP TS 38.331 UAC framework.



Fundamental causality constraint

ACB throttles only after observing $Q_{\text{real}} > 0.8K$. During the build-up, requests are silently dropped.

Empirical evidence

Seo et al. (2021): Random access delay requirements can be violated when the RA procedure enters unstable or saturated operating regions.

This supports the claim that reactive ACB/UAC is too late under bursty overload: once congestion is observed, delay and drops may already be unavoidable.

Why existing approaches fall short



Reactive ACB / UAC

Strength

Simple, standards-compliant

Limitation

Reacts after saturation; inherent control-loop latency



ML / DNN predictors

Strength

High accuracy on stationary traffic

Limitation

$O(W)$ inference, $W \sim 10^5 - 10^7$; ms latency — impractical on gNB hardware



Standard Digital Twin

Strength

Proactive control possible

Limitation

Model drift under non-stationary traffic; mirroring error grows unbounded



Gap: no DT approach suppresses model drift via an online calibration loop during the attack interval. This work closes that gap.

MMPP/M/c/K queue

Two-state CTMC: Markov-Modulated Poisson Process



Generator matrix

$$Q = \begin{bmatrix} -\alpha & \alpha \\ \beta & -\beta \end{bmatrix} ; \quad \text{sojourn} \sim \text{Exp}(|Q_{ii}|)$$

Exact CTMC dynamics in DES — no matrix exponential needed.


Physical gNB RACH

Parameter	Value
Preamble channels	$c = 54$
Buffer capacity	$K = 200$
Service time	$\text{Exp}(1)$



Stationary blocking governed by classical M/c/K formulas;
non-stationarity comes from MMPP modulation.

Calibrated Digital Twin architecture



EMA predictor and ACB control policy



EMA arrival-rate predictor

$$\text{EMA}_t = \alpha \cdot \lambda_{\text{cur}} + (1 - \alpha) \cdot \text{EMA}_{\{t-1\}}$$

$$\lambda_{\text{cur}} = n_{\text{win}} / \Delta$$

- Initial weight $\alpha_0 = 0.20$
- α is dynamic — updated by the calibration loop (Stage 3)

Key insight: p_{acb} rises before Q_{real} reaches K — proactive throttling, not reactive.



Stochastic barring probability

$$p_{\text{acb}} = \max(0, \min(1, 1 - (K - Q_{\text{real}}) / \max(\text{EMA}_t \cdot \Delta, 1e-9)))$$

- Minimum probability to prevent $Q_{\text{real}} \geq K$ next step
- Applied as a Bernoulli trial per request
- Maps to uac-BarringInfo in SIB1 (3GPP TS 38.331)

Continuous calibration loop

State-mirroring error

$$E(t) = | Q_{\text{real}}(t) - Q_{\text{virtual}}(t) |$$

Asymmetric step rule

```
if E > ε :  
    α ← min(α_max, α + δ+)  
else:  
    α ← max(α_min, α - δ-)
```

ε**10**

calibration threshold

δ⁺**0.05**

fast attack response

δ⁻**0.01**

slow decay (stability)

Attack onset: δ⁺ = 0.05 (fast)

- EMA becomes high-frequency
- Outdated observations discarded quickly

Normal traffic: δ⁻ = 0.01 (slow)

- Statistical stability preserved
- Avoids oscillation around ε

Calibrated DT control loop

 Every Δ control step

```

1.  $\lambda_{\text{cur}} \leftarrow n_{\text{win}} / \Delta$ 
2.  $\text{EMA}_t \leftarrow \alpha \cdot \lambda_{\text{cur}} + (1 - \alpha) \cdot \text{EMA}_{\{t-1\}}$ 
3.  $d \leftarrow \max(\text{EMA}_t \cdot \Delta, 1e-9)$ 
4.  $p_{\text{acb}} \leftarrow \max(0, \min(1, 1 - (K - Q_{\text{real}}) / d))$ 
5.  $Q_{\text{virtual}} \leftarrow \text{clip}(Q_{\text{virtual}} + \text{EMA}_t \cdot (1 - p_{\text{acb}}) \cdot \Delta - c \cdot \Delta, 0, K)$ 
6.  $E \leftarrow |Q_{\text{real}} - Q_{\text{virtual}}|$ 
7. if  $E > \varepsilon$  :  $\alpha \leftarrow \min(\alpha_{\text{max}}, \alpha + \delta^+)$ 
8. else      :  $\alpha \leftarrow \max(\alpha_{\text{min}}, \alpha - \delta^-)$ 
9.  $\text{drop\_rate} \leftarrow p_{\text{acb}}$ 

```

Complexity

Calibrated DT

$O(1)$

Fixed arithmetic · 4 state variables · μs control budget

DNN predictors

$O(W)$

$W \sim 10^5 - 10^7 \cdot \text{ms latency} \cdot \text{impractical on gNB}$

State: $\{ \text{EMA}_t, Q_{\text{virtual}}, \alpha, n_{\text{win}} \}$

Simulation setup



Experimental configuration

- Discrete-event simulation: SimPy, Python 3.10
- $T_{sim} = 1000$ s per run
- $N = 10$ independent seeds (Monte Carlo)
- 95 % CI: $\bar{x} \pm 1.96 \cdot \hat{\sigma} / \sqrt{N}$

Three methods compared

- Reactive Baseline

threshold 0.8 K, drop 0.5
- DT no-calibration


α fixed
- Calibrated DT

asymmetric step rule

Key parameters

Parameter	Symbol	Value
Preamble channels	c	54
Buffer capacity	K	200
Normal arrival rate	λ_{low}	10 req/s
Burst arrival rate	λ_{high}	500 req/s
Control interval	Δ	1 s
Calibration threshold	ϵ	10
Initial EMA weight	α_0	0.20

State-mirroring error — 42% reduction




42 %
reduction in E_{mean}

E_{mean} (N = 10, 95 % CI)	
DT (no calib.)	85.03 $\pm$ 4.40
Calibrated DT	49.14 $\pm$ 3.09

Takeaway

Calibration suppresses model drift continuously during the burst — CIs do not overlap.

Queue stability



P_{success} (N = 10, 95 % CI)

Reactive	0.310 $\pm$ 0.015
DT (no calib.)	0.301 $\pm$ 0.011
Calibrated DT	0.309 $\pm$ 0.019




statistically equivalent to baseline; no throughput penalty

Interpretation

- Reactive only acts after $Q > 0.8 K$
- Calibrated DT raises p_{acb} before saturation

Queueing delay — CDF



CDF left-shift

Calibrated DT curve sits left of the others — lower delay across all reported percentiles.

Tail suppressed




P90 / P99 queueing delay roughly halved — critical for URLLC latency budgets.

Mechanism

- Reactive admits requests into a deeply loaded queue
- Calibrated DT keeps Q_{real} low → short waits

ϵ -sensitivity analysis

Epsilon Sensitivity Analysis (Calibrated DT, seed=42)



Tested values

$\epsilon \in \{ 5, 10, 20, 40, 80 \}$

Single-seed sweep (seed = 42)

Interpretation

- $\epsilon = 5 \rightarrow$ over-calibrates on normal traffic
- $\epsilon = 80 \rightarrow$ calibration rarely triggers
- $\epsilon = 40 \rightarrow$ best composite in this sweep


Adopted default

$\epsilon = 10$

conservative — just above the normal-traffic noise floor

$O(1)$ complexity and O-RAN integration

O-RAN integration



Deployment notes

- $\Delta = 1$ s fits Near-RT RIC range (10 ms – 1 s)
- No protocol modifications required
- Vendor-independent xApp — standard E2 service models
- $O(1)$ per step — fixed-cost arithmetic, suitable for tight gNB control budgets after hardware benchmarking

Calibrated DT runs in μ s vs. ms-class DNN inference — practical for commercial gNB software.

Findings and future work

42 %

reduction in state-mirroring error vs. uncalibrated DT

$\approx$

P_success statistically equivalent to baseline; P90/P99 delay reduced

$O(1)$

deployable in commercial gNB software via O-RAN xApp



Proactive control in non-stationary 5G environments requires both a sound control policy and tight virtual-physical alignment.

Multi-cell extension

Current model treats one gNB cell independently. Real deployments involve handover correlations and inter-cell interference — a distributed DT with inter-cell state sharing is required.

RL-based calibration

Fixed δ^+ and δ^- are vulnerable to slow ramp-up attacks. An RL agent can learn step sizes dynamically using state-mirroring error E as a reward signal.



Questions?

5G RRC Signaling Storm Mitigation via a Proactive Digital Twin Architecture with Continuous Calibration

Ömer Dik Yol

dik yol25@itu.edu.tr

BLG 632E · ITU

Selected references

- [b1] A. Tabiban, H. A. Alameddine, M. A. Salahuddin, and R. Boutaba, "Signaling Storm in O-RAN: Challenges and Research Opportunities," IEEE Communications Magazine, vol. 62, no. 6, pp. 58-64, Jun. 2024. doi: 10.1109/MCOM.002.2300317.
- [b2] B. H. Lee, H. S. Lee, S. Moon, and J. W. Lee, "Enhanced Random Access for Massive Machine Type Communications," IEEE Internet of Things Journal, vol. 8, no. 8, pp. 7046-7064, Apr. 2021. doi: 10.1109/JIOT.2020.3038148.
- [b3] M. Grieves and J. Vickers, "Digital Twin: Mitigating Unpredictable, Undesirable Emergent Behavior in Complex Systems," in Transdisciplinary Perspectives on Complex Systems, Springer, 2017, pp. 85-113. doi: 10.1007/978-3-319-38756-7_4.
- [b4] ETSI, "5G; NR; Radio Resource Control (RRC); Protocol specification," ETSI TS 138 331 V17.3.0, 3GPP TS 38.331 version 17.3.0 Release 17, Jan. 2023.
- [b5] L. Miuccio, D. Panno, and S. Riolo, "Joint Control of Random Access and Dynamic Uplink Resource Dimensioning for Massive MTC in 5G NR Based on SCMA," IEEE Internet of Things Journal, vol. 7, no. 6, pp. 5042-5063, Jun. 2020. doi: 10.1109/JIOT.2020.2974402.
- [b6] J.-B. Seo, W. T. Toor, and H. Jin, "Analysis of Two-Step Random Access Procedure for Cellular Ultra-Reliable Low Latency Communications," IEEE Access, vol. 9, pp. 5972-5985, Jan. 2021. doi: 10.1109/ACCESS.2020.3048824.
- [b7] S. K. Sharma and X. Wang, "Toward Massive Machine Type Communications in Ultra-Dense Cellular IoT Networks: Current Issues and Machine Learning-Assisted Solutions," IEEE Communications Surveys & Tutorials, vol. 22, no. 1, pp. 426-471, 2020. doi: 10.1109/COMST.2019.2916177.
- [b8] R. Dong, C. She, W. Hardjawana, Y. Li, and B. Vucetic, "Deep Learning for Radio Resource Allocation With Diverse Quality-of-Service Requirements in 5G," IEEE Transactions on Wireless Communications, vol. 20, no. 4, pp. 2309-2324, Apr. 2021. doi: 10.1109/TWC.2020.3041319.
- [b9] C. Zhang, P. Patras, and H. Haddadi, "Deep Learning in Mobile and Wireless Networking: A Survey," IEEE Communications Surveys & Tutorials, vol. 21, no. 3, pp. 2224-2287, 2019. doi: 10.1109/COMST.2019.2904897.
- [b10] H. X. Nguyen, R. Trestian, D. To, and M. Tatipamula, "Digital Twin for 5G and Beyond," IEEE Communications Magazine, vol. 59, no. 2, pp. 10-15, Feb. 2021. doi: 10.1109/MCOM.001.2000343.
- [b14] M. Polese, L. Bonati, S. D'Oro, S. Basagni, and T. Melodia, "Understanding O-RAN: Architecture, Interfaces, Algorithms, Security, and Research Challenges," IEEE Communications Surveys & Tutorials, vol. 25, no. 2, pp. 1376-1411, 2023. doi: 10.1109/COMST.2023.3239220.